Transformers and LLMs

Transformers

The transformer architecture introduced self-attention and replaced recurrent neural networks for many NLP tasks.

Self Attention

Self-attention lets every token attend to every other token in a sequence.

LLMs

Large Language Models are trained on massive datasets and can generate human-like text.

Fine-tuning

Models can be fine-tuned for specific tasks using domain-specific data.

Inference

During inference, prompts are converted to tokens, processed by the model, and decoded into text.